

THE EFFECTS OF PRIVACY REGULATION ON THE MARKET FOR STOLEN DATA

Anderson Frailey *

August 17, 2026

[Latest version available here](#)

Abstract

I propose a model of the stolen data economy to show that a privacy regulation decreases the equilibrium quantity of stolen data bundles available and level of scamming. The latter effect alleviates the market's asymmetric information problem, increasing breach probability for the most data-rich targets. Using a novel dataset of data breaches, I show the General Data Protection Regulation caused a 60 percent reduction in the number of data breaches in the stolen data market, no corresponding change in the aggregate amount of data available, and a 63 percent increase in the amount of data the average individual breach contains. (JEL: D82 K24, L51, L86)

*University of Maryland, Baltimore County, Email: afrailey@umbc.edu. UMBC Department of Economics 1000 Hilltop Circle, Public Policy Building, 3rd Floor, Baltimore, MD 21250. I would like to thank my advisors at the University of Virginia—Amalia Miller, Lee Lockwood, and Denis Nekipelov—for their guidance and support on this paper. Conversations with Sarah Turner, Daniel Kwiatkowski, Danielle Citron, Diego J. Jiménez Hernández, and Mike Andrews were incredibly helpful in this paper's development. I received valuable comments and feedback from participants in the Public and Labor Student Workshop and Economics Research Colloquium at the University of Virginia, as well as the 2024 AEA Summer Mentoring Pipeline Conference. Finally, this paper would not have been possible without the data provided by SpyCloud. For that I thank Pablo Maceda, Ronak Patel, Wallis Romzek, and Trevor Hilligoss. The conclusions and views expressed in this paper do not reflect those of SpyCloud or any persons affiliated with SpyCloud. All remaining mistakes are my own.

1 Introduction

Personal data collected by businesses, schools, and most other modern organizations is frequently exposed in data breaches. In 2025, researchers found 2.14 billion national identifiers, 90 million credit cards, and 13 million passports accessible online (SpyCloud, 2026). Reducing the number of data breaches and amount of data leaked is one of the goals of data privacy regulations, but whether these regulations are successful in that objective has yet to be thoroughly examined. This is in part because it is a difficult question to answer causally. Privacy regulations typically include data breach disclosure requirements, increasing the probability a breach is observed.

This paper contributes to solving this problem in three ways. First, I propose studying how the European Union’s General Data Protection Regulation (GDPR), one of the most comprehensive data privacy laws in the world, has affected the market for stolen data. In the stolen data market, bundles of data taken in data breaches are traded for money, bragging rights, and reputation. Focusing on the stolen data market, rather than other measures of the number of data breaches, circumvents the observation problem because the number of postings in the market will be affected by the provisions of the regulation designed to reduce the frequency of data breaches, those that reduce data collection and increase cybersecurity requirements, but not the disclosure requirements.

Next, I propose a model of the stolen data economy to show that, although privacy regulation will likely decrease the expected number of data breaches, it may not reduce the amount of data available in the stolen data market. In fact, such a regulation could even increase the amount of stolen data available. This is because a privacy regulation reduces the equilibrium level of fraudulent data sellers, alleviating the market’s asymmetric information problem. More generally, the model provides an example of how regulating a legal market can influence an illicit market that is not directly touched by the regulation.

Finally, I introduce a novel dataset of data packages traded on the stolen data market to estimate how the GDPR changed aggregate—number of data breaches and total amount of data available—and product-level—individual breach size and contents—outcomes in the stolen data market. To the best of my knowledge, this is the first use of such data in the economics literature.

A privacy regulation is a supply shock in the stolen data market. By reducing the amount of data collected (Demirer et al., 2024) and imposing cybersecurity requirements, the regulation makes legally collected data, the market’s key input, more difficult to obtain. However, modeling the effect of a privacy regulation is not as simple as modeling a negative supply shock because of the stolen data market’s asymmetric information problem. In stolen data markets, participants are at least semi-anonymous and only sellers can fully observe the true value of their goods prior to purchase, making buyers willing to pay less for data than it is actually worth (Akerlof, 1970; Herley and Florêncio, 2010). The model I propose treats the level of scamming in the stolen data market as an equilibrium object influenced by the regulation. The key theoretical finding is that, under the right conditions, even if privacy regulations like the GDPR reduce the number of data breaches, they may not reduce the amount of data stolen. This is because the desired reduction in

breach probability is concentrated among mid-sized data collectors, while the most data-intensive organizations see their breach probability increase. The driving force behind the latter effect is a decrease in the equilibrium level of fraudulent data in the stolen data market. Effectively, a privacy regulation alleviates the lemons problem in the illicit market.

Empirically, I employ a unique dataset of stolen data packages posted in the market. Each observation is of an individual data package and contains information on the organization the data originated from, as well as the amount and types of data included. It is important to note that these data only cover what is available online, not necessarily everything that was stolen in a given data breach. The two may differ if a hacker decides to keep some data for themselves or that some of the data is not worth selling. Each package is labeled as being available before or after the GDPR, and whether the data it contains should have been protected by the GDPR.

I combine the individual data packages to create a country-quarter level panel spanning January 2017 to November 2023 that tracks the number of data breaches and records available in stolen data markets that originate in a given country. I use this panel to estimate a difference-in-differences model measuring the effect of the GDPR on those two outcomes. I also break the post-GDPR period into short-run and long-run periods to measure if and how the effects changed over time.

I find that the GDPR caused the number of data breaches originating from regulated countries to decrease by approximately 60 percent overall, with the long-run decrease being slightly larger than the short-run (61 percent versus 55 percent). Despite this, I find no statistically significant change in the total amount of data available. The granularity of my data allows me to estimate how the composition of the individual data packages changes to explain the lack of change in number of records.

At the individual data package level, I estimate how the content of the data packages—the number of records, amount of personally identifiable information (PII), and number of unique types of data—changes after the GDPR. As with the aggregate effects, I estimate both an overall change and separate short- and long-run changes. I find that the size of data packages, in terms of number of records, originating in regulated countries increases 68 percent in the long-run, while there is no statistically significant short-run change. The fraction of those records that are considered PII and number of unique data types in these packages do not change in any measured time period.

The increase in size of the data packages explains how the number of data packages available could fall without an accompanying decrease in the number of records. Additional empirical evidence of a shift towards larger targets is in the UK Cyber Security Breach Survey, which shows that small organizations (those with fewer than 50 employees) make up 95 percent of reported breaches in the 2017 survey, but only 48 percent in the 2022 survey. Large organizations (those with 250 or more employees) increase their share from less than one percent to approximately 24 percent (Department for Digital, Culture, Media and Sport, 2022). These results are consistent with the breach composition change predicted by the model.

The effects of data privacy and security legislation have been studied in a number of contexts including healthcare (A. Miller and Tucker, 2009; A. Miller and Tucker, 2011; A. Miller and Tucker,

2018) and online advertising (Goldfarb and Tucker, 2011). The GDPR-specific literature covers its effects on firm performance (Koski and Valmari, 2020; C. Chen, Frey, and Presidente, 2022; Goldberg, Johnson, and Shriver, 2024), competition (Johnson, Shriver, and Goldberg, 2023), investment (J. Jia, Jin, and Wagman, 2021; Kircher and Foerderer, 2021; Janßen et al., 2022; Jia et al., 2025), and data collection (Aridor, Che, and Salz, 2021; Lukic, K. M. Miller, and Skiera, 2023; Demirer et al., 2024). These papers typically find that the regulation had negative effects on business operations—decreases in competition, investment, and firm performance—or changes whose net welfare effects are more ambiguous, such as decreased data storage. Jones and Tonetti (2020) provide a theoretical backing to these results, showing that government policies restricting data use can increase privacy at the cost of reducing efficiency.

While I do not attempt to calculate the overall welfare effects of the policy, this paper does find a seemingly unambiguous benefit of the GDPR: the reduction in the number of data packages online. But even with this reduction, the extent to which individuals benefit is unclear given that there was no accompanying reduction in the number of records available. It is possible that, while there are fewer breaches, those that remain contain enough information to leave the affected individuals no better off than before.

That said, there is also evidence that privacy regulations may be beneficial in other ways. Economists have shown that the cost to a firm of suffering a data breach is much lower than the cost of other incidents—such as fraud or corruption—when not including other social and consumer costs (Romanosky, 2016). If organizations are underinvesting in security relative to the social optimum because they do not fully internalize the cost of a data breach, imposing security requirements is one approach to resolving that externality problem. As might the data breach disclosure provisions of these regulations. For example, Romanosky, Telang, and Acquisti (2011) show that data breach disclosure laws reduce identity theft stemming from data breaches by 6.1 percent.

Significant work studying stolen data markets has been conducted by criminologists, who have derived some estimates of their sizes and products offered (Franklin et al., 2007; Holt and Lampke, 2010; Holt, Smirnova, and Chua, 2016). These papers conduct in-depth, descriptive studies of a handful of individual forums where data is sold. They do not study how public policy can have trickle-down effects on these markets. My unique dataset allows me to fill that gap in this area of the criminology literature, and extend the contribution into the economics of crime.

The model I present conceptually aligns with Becker (1968). Hackers decide whether to attempt a data breach based on the perceived costs and benefits of doing so. When the costs increase and the benefits decrease, there are fewer breaches. The costs and benefits are not perfectly observable, requiring hackers to base their actions on their beliefs of data collection practices and how well potential targets have protected themselves. This is similar to the mechanisms in Ayres and Levitt (1998) and Braakmann, Chevalier, and Wilson (2024). In Ayres and Levitt, car thieves could not observe which vehicles had tracking devices installed, but were aware of which areas had higher installation rates. The higher likelihood that a stolen car could be tracked caused them to steal

fewer cars in those areas. In Braakmann, Chevalier, and Wilson, an increase in the price of gold motivated burglars to target homes in areas where homeowners were expected to store more gold. By causing a reduction in the expected value (akin to the Braakmann, Chevalier, and Wilson mechanism in reverse) and an increase in the expected cost (à la Ayres and Levitt) of breaching European organizations, the regulation incentivizes hackers to change who they target.

The remainder of this paper is structured as follows. Section 2 presents a brief background of the GDPR and its relationship to the stolen data market. Section 3 formally presents the model of the stolen data market discussed earlier. In Section 4 I describe the data used in this study. My empirical strategy is defined in Section 5 and the results are presented in Section 6. I offer concluding remarks on the implications and limitations of this study, and paths for future research in Section 7.

2 Background

The GDPR is a broad-reaching regulation that governs the collection and processing of personal data by covered organizations. It explicitly states under what circumstances data collection is considered lawful, and prohibits the processing of sensitive data with few exceptions. Additionally, the GDPR gives individuals the right to have their data deleted, transferred, or rectified; requires detailed record keeping on data collection, impact assessments prior to data processing, and the designation of a Data Protection Officer; and increases cybersecurity investment requirements. Notification requirements and large fines also significantly increase the cost of suffering a data breach. Demirer et al. (2024) estimate the GDPR increased the cost of data storage by 20 percent on average, and caused a 26 percent decrease in data storage among firms in the EU relative to comparable American organizations.

By reducing the amount of data collected and requiring increased cybersecurity, the GDPR reduces the availability of the primary input for illicit data markets: data. In that sense, the GDPR is a negative supply shock for data thieves. Breaking down the stolen data supply chain into two parts—legal data collection and data theft—makes this effect clear.

Legal data collectors gather information on their customers, employees, and users for marketing and general operations. These agents are directly regulated by the GDPR. They change their data collection and security practices to comply with the law.

After its legal collection, data is stolen by cyber criminals. For hackers, each potential target has an expected value and cost of hacking. They will choose how to allocate their efforts across targets based on these expectations and the conditions of the market. By reducing the amount of data collected and requiring organizations to invest in security, privacy regulations should decrease the value and increase the cost of breaching regulated organizations. This changes the set of targets worth hacking, and the supply available in the market.

What makes this problem more complicated than a simple negative supply shock is the presence of scammers in the stolen data market. Goods sold in these illicit markets tend to sell at lower

prices than they are worth because buyers know they can get scammed (Herley and Florêncio, 2010). Treating the amount of scamming present in the market as part of its equilibrium adds nuance to the predicted effects of a privacy regulation, as I show in the following section.

3 A Model of the Stolen Data Economy




FIGURE 1
The Stolen Data Supply Chain

Figure 1 depicts how stolen data is produced and used. Legal data collectors—companies, schools, governments, and similar entities that hold customer, user, and employee data—initially decide what and how much information to collect. Collecting data comes at a cost. They must pay to gather it, keep it stored, and respond to user requests regarding their data. There is also the persistent risk of suffering a data breach and the additional costs that result: sending notifications to those affected, offering credit monitoring, performing security audits, legal costs, and fines imposed by governments. Organizations invest in cybersecurity to mitigate this risk. Prevention methods can be technical, consistently patching software vulnerabilities and encrypting data, or non-technical, teaching employees to detect phishing emails and improving their password management. The goal of these measures is to make it more difficult for data to fall into the wrong hands.¹

In the second stage, hackers choose which targets to breach. All else equal, targets who collect more (and more valuable) data are more attractive to hackers. Investing more in security makes targets less attractive.

Successfully stolen data is traded in online black markets found on platforms such as Telegram or data breach sites. Hackers, the market’s suppliers, advertise their products by describing what

¹This goal is not always achieved. A. Miller and Tucker (2011) find that use of encryption technology is actually associated with an increase in reports of data loss. They speculate this is due to firms and employees being less careful with the data they hold after it is encrypted.

is in the data package and where it originated from (Holt and Lampke, 2010). Because they are anonymous, online, and illegal, these markets are especially vulnerable to adverse selection and asymmetric information problems. Many platforms have introduced reputation systems to solve such problems and facilitate trade.² However, most transactions are still one-offs, limiting the effectiveness of these systems (Vu et al., 2020). A consequence of these information problems is that goods tend to sell for significantly less than their true value as buyers account for the risk of being scammed (Herley and Florêncio, 2010).

Privacy regulations are a negative supply shock in the market along two dimensions. First, they reduce the amount of data stored by organizations (Demirer et al., 2024), decreasing the benefit of breaching regulated targets. Second, they typically require increased investment in cybersecurity, increasing the cost of breaching regulated targets. Data from the United Kingdom Cyber Security Breaches Survey shows that nearly two-thirds of respondents made operational changes in response to the GDPR. Among those that made changes, 100 percent reported making changes related to the cybersecurity policies and practices (table 1). Taken as a whole, privacy regulations increase the cost of acquiring the key input to the stolen data market: the data itself. This changes the incentives of hackers. The amount of data that can be extracted from a regulated organization decreases, while the marginal cost of breaching one increases, encouraging changes in the optimal effort allocation.

Table 1: Percentage of Organizations Reporting Operational Changes Due to the GDPR

Survey Year	Any Change		Change in Cybersecurity	
	2018	2019	2018	2019
Overall	12.75%	63.71%	100.00%	100.00%
Small	9.91%	61.56%	100.00%	100.00%
Medium	28.71%	90.66%	100.00%	100.00%
Large	52.91%	95.81%	100.00%	100.00%

Note: Department for Digital, Culture, Media and Sport (2022), author’s calculations. Respondents were asked “Has your organisation made any changes or not to the way you operate in response to GDPR?” and “Have any of these changes been related to your cyber security policies or processes, or not?” The fraction of respondents answering yes to the first question is in the first two columns. The fraction answering yes to the second question, among those answering yes to the first, is in the last two columns.

To formalize the above, I propose a model of the stolen data economy with two parts: the stolen data market and legal data collectors. Buyers and sellers in the stolen data market trade data packages—bundles of data advertised as being stolen from legal data collectors. Legal data collectors choose how much data to collect and their security investment levels. The latter’s choices will be influenced by equilibrium in the stolen data market.

²Some trading platforms have also created contract systems that set expectations for the parties involved in a transaction and help build suppliers’ reputation (Vu et al., 2020). Often, suppliers will give away their stolen data rather than sell it to help build their reputation.

3.1 The Stolen Data Market

Buyers in the stolen data market know that the data packages they purchase are either legitimate and worth V , or sold by a scammer and worth nothing. They have no way of distinguishing between legitimate and fake data prior to purchase. Based on their knowledge of the market, they form beliefs over the probability a data package is legitimate. Denoting the portion of legitimate data packages in the market q , the expected value of a bundle of data is

$$\mathbb{E}[V] = qV + (1 - q)0 = qV$$

Buyers have a skill parameter κ that determines how much benefit they receive from a given data package. For a bundle of data purchased at price p , the expected net utility received is

$$U_b = \kappa qV - p$$

For simplicity, assume $\kappa \sim \text{Uniform}[0, 1]$ and that buyers will only buy a bundle of data if the expected utility is non-negative, the market demand from a unit mass of buyers is

$$Q^D(p, q, V) = 1 - \frac{p}{qV} \quad (1)$$

Legitimate sellers are heterogeneous in their ability, γ , which determines their cost of hacking an organization given its security investment (S). The profit earned by a legitimate seller on the sale of a data package is:

$$U_h = p - \gamma S$$

Again for simplicity, assume $\gamma \sim \text{Uniform}[0, 1]$. The portion of hackers that can earn profit from, and therefore want to breach, a target is $\frac{p}{S}$. With a unit mass of legitimate sellers, total legitimate supply will be

$$Q_\ell^S(p, S) = \frac{p}{S}$$

Scammers are included in the model via a single representative scammer that can produce fake data at a cost, which I normalize to zero.³ The scammer is aware that the more fake data they create and attempt to sell, the less buyers are willing to pay. Factoring in this effect, they will choose the optimal amount of fake data to create, effectively choosing q . The scammer's profit maximization problem is

$$\max_q (1 - q) \left(1 - \frac{p}{qV} \right) p$$

³I am assuming that the cost of generating fake data is fixed and does not depend on the value of data they are impersonating. Allowing the cost of fake data generation to vary with data quality would complicate the problem without adding to the intuition.

Solving this yields the optimal q given the market price and value of the data being faked:

$$q = \sqrt{\frac{p}{V}} \quad (2)$$

Defining Γ as the quantity of fake data bundles the scammer puts into the market, total market supply is

$$Q^S(p, S) = Q_\ell^S(p, S) + \Gamma$$

Equilibrium in the stolen data market consists of two conditions:

1. Buyer beliefs over q are accurate i.e.,

$$q = \frac{Q_\ell^S}{Q^S} = \frac{p/S}{p/S + \Gamma} \quad (3)$$

2. The market clears: $Q^S = Q^D$.

Under the second condition, the denominator of equation 3 can be replaced by equation 1 and the result rearranged to isolate q :

$$q = \frac{p}{V} + \frac{p}{S} \quad (4)$$

By the first equilibrium condition, the equilibrium price as a function of the target's security and value is found by setting equations 2 and 4 equal:

$$p^* = \frac{S^2 V}{(S + V)^2} \quad (5)$$

This price is subsequently used to find equilibrium q^* :

$$q^* = \frac{S}{S + V}$$

3.2 Legal Data Collectors—Pre-GDPR

The stolen data market's equilibrium determines the probability a legal data collector is breached. Given their choice of data collection and security, the probability a legal data collector is breached is equal to the fraction of legitimate sellers that find it profitable to hack them. Plugging the market price given by equation 5 into the legitimate supply curve, breach probability, $\pi(S, V)$, is

$$\pi(S, V) = \frac{p^*}{S} = \frac{SV}{(S + V)^2} \quad (6)$$

If breached, legal data collectors will incur a loss of

$$L(S, V) = \frac{\Lambda V^\psi}{\delta S}, \quad 0 < \psi < 2$$

Considering the costs and benefits of data collection, the legal data collector optimization problem is

$$\max_{S,V} AV^\phi - \pi(S,V)L(S,V) - C(V) - C_S(S), \quad 0 < \phi < 1$$

where $C(V)$ and $C_S(S)$ are the cost of data collection and security, respectively, and A is a productivity term.

Assume that the value of data collection takes the functional form:

$$V = (\alpha_1 d_1^\rho + \dots + \alpha_J d_J^\rho)^{1/\rho}$$

where d_i is the amount of data of type i collected. To simplify the optimization, define the security-to-data value ratio as $r = \frac{S}{V}$. Finally, define the two cost functions as:

$$C(V) = \sum_{j=1}^J \omega_j d_j, \quad C_S(S) = \omega_S S = \omega_S r V$$

Using these definitions, substituting in $S = rV$ for S , and reducing the problem to using only two data types for simplicity, the optimization problem becomes

$$\max_{r,d_1,d_2} A (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\phi/\rho} - \frac{\Lambda}{\delta (1+r)^2} (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\frac{\psi-1}{\rho}} - \sum_{j=1}^2 \omega_j d_j - \omega_S r (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{1/\rho}$$

Taking the first order conditions and simplifying yields the conditional factor demands for d_1 , d_2 , and r :⁴

$$d_1(V) = \frac{V}{B_1^{1/\rho}}, \quad d_2(V) = \frac{V}{B_2^{1/\rho}},$$

$$r(V) = \left[\frac{2\Lambda}{\delta \omega_S} \right]^{\frac{1}{3}} V^{\frac{\psi-2}{3}} - 1$$

The legal data collector's optimal V is found by setting the marginal benefit of V equal to its marginal cost:⁵

$$\underbrace{\phi AV^{\phi-1}}_{MB} = \underbrace{\frac{\omega_1}{B_1^{1/\rho}} + \frac{\omega_2}{B_2^{1/\rho}} + \left(\omega_S \left[\frac{2\Lambda}{\delta \omega_S} \right]^{1/3} + \frac{\Lambda}{\delta \left[\frac{2\Lambda}{\delta \omega_S} \right]^{2/3}} \right) \frac{\psi+1}{3} V^{\frac{\psi-2}{3}} - \omega_S}_{MC} \quad (7)$$

⁴Derivations are in section A of the appendix. The B_i terms are:

$$B_1 = \left[\alpha_2 \left[\frac{\omega_2}{\omega_1} \frac{\alpha_1}{\alpha_2} \right]^{\frac{\rho}{\rho-1}} + \alpha_1 \right], \quad B_2 = \left[\alpha_1 \left[\frac{\omega_1}{\omega_2} \frac{\alpha_2}{\alpha_1} \right]^{\frac{\rho}{\rho-1}} + \alpha_2 \right]$$

⁵Derivations are once again in section A of the appendix

This does not have a closed form solution, but one is not necessary to compute the relevant comparative statics in section 3.4.

3.3 Legal Data Collectors—Post-GDPR

For the purpose of this paper, the GDPR had two main effects: increasing the cost of data collection (ω_1 , ω_2 , and Λ) and adding a minimum security constraint that is a function of the value of data collected.⁶ Modeling the latter as a minimum security-to-data value ratio constraint, $\underline{r}$, the post-GDPR legal data collector optimization problem is:

$$\begin{aligned} \max_{r, d_1, d_2} & A (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\phi/\rho} - \frac{\Lambda}{\delta (1+r)^2} (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\frac{\psi-1}{\rho}} - \sum_{j=1}^2 \omega_j d_j - \omega_S r (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{1/\rho} \\ \text{s.t. } & r \geq \underline{r} \end{aligned}$$

The conditional factor demands for d_1 and d_2 are unchanged, as is the conditional factor demand for r when the constraint does not bind. When the constraint does bind, $r = \underline{r}$, which changes the marginal cost of V . The constrained legal data collectors' new V^* solves

$$\underbrace{\phi A V^{\phi-1}}_{MB} = \underbrace{\frac{\omega_1}{B_1^{1/\rho}} + \frac{\omega_2}{B_2^{1/\rho}} + (\psi-1) \frac{\Lambda}{\delta (1+\underline{r})^2} V^{\psi-2} + \omega_S \underline{r}}_{MC}$$

This also only has a closed form solution under a specific value of ϕ , which I will again decline to assume.

3.4 Comparative Statics: The GDPR Effect

There are three GDPR-driven changes of interest: the change in the number of breaches, the change in average breach size (in terms of the amount of data), and the change in the total amount of data available in the market. These will each be affected by both the increase in data collection costs and the security constraint.

Focusing first on the security constraint, two features of the legal data collector's optimization solution are important for determining the GDPR's effect: V^* is monotonically increasing in A , and r^* is monotonically decreasing in A .⁷

Substituting $S = rV$ into the breach probability function,

$$\pi = \frac{SV}{(S+V)^2} = \frac{r}{(1+r)^2}$$

⁶The increasing cost of data collection is a product of the GDPR's additional obligations for data collectors, such as requiring impact assessments prior to data processing and sending data breach notifications. The minimum security requirement is motivated by Article 32, which requires organizations "implement appropriate technical and organisational measures to ensure a level of security appropriate to the risk" and lists four specific measures.

⁷Derivations are again in section A of the appendix.




FIGURE 2
 $\pi(r)$: Breach Probability as a Function of r

This is non-monotonic in r (see figure 2) and, therefore, A . As r increases (and A decreases), π initially increases before reaching its maximum and falling.

Defining $A = A^\dagger$ as the value of A where $r(A^\dagger) = \underline{r}$, all targets with $A < A^\dagger$ will be unaffected by the constraints.⁸ All others will need to increase their security ratio to $\underline{r}$, giving each constrained target a breach probability of $\pi(\underline{r}) \equiv \bar{\pi}$. A convenient property of the π function is that $\pi(r) = \pi(\frac{1}{r})$. Using this, define $A^{\dagger\dagger}$ as the value of A such that $\pi(r(A^{\dagger\dagger})) = \pi(\frac{1}{\underline{r}})$. If $\underline{r} > 1$, the security constraint will reduce breach probability for targets with $A \in (A^\dagger, A^{\dagger\dagger})$, but *increase* breach probability for those with $A > A^{\dagger\dagger}$ (see figure 2).

The overall effect of the GDPR on the expected number of breaches therefore depends on A 's distribution, $F(A)$, and which of these changes dominates. If a sufficiently small mass of legal data collectors falls in the $A > A^{\dagger\dagger}$ group, the expected number of breaches will fall. Regardless of the exact distribution of A , this changes in the composition of breaches: higher A targets will make up a larger portion of the post-GDPR breaches than pre-GDPR breaches. These are organizations with high levels of data collection.

The expected amount of data in the market depends on both π and the total amount of data collected $D(A) = d_1(A) + d_2(A)$:

$$\mathbb{E}[\text{Data}] = \int \pi(A) D(A) dF(A)$$

While the effect of the security constraint on π is ambiguous, it is not for total data collection. Imposing higher security requirements increases the marginal cost of data collection, reducing optimal $D(A)$. Increasing the cost of data collection will similarly reduce $D(A)$.

The GDPR's total effect on the amount of data in the market depends on the relative sizes of the aforementioned compositional change and the reduction in data collection. The shift towards

⁸To the extent organization size correlates with amount of data held, this partitioning explains why only 61 percent of small organizations in the UK Cyber Security Breaches Survey reported making operational changes after the GDPR, compared to 90 and 95 percent of medium and large organizations (table 1).




FIGURE 3
 $q(r)$: Portion of Legitimate Data Packages as a Function of r

more data-rich targets increases the expected amount of data, while the decrease in data collection by all targets reduces the expected data quantity. The net change in total data on the market will depend on which of these effects dominates, but the latter has some nuance. Each type of target reduces its data collection, including the now more prominent high A types. But if their reduction in data collection is small enough, they will still have more data than the mid- A types had prior to the GDPR. If this is the case, the composition effect still increases the expected amount of data stolen. The same logic applies to determining how the expected breach size changes. If the composition effect dominates, the average amount of data in a breach increases.

To understand the driving force behind the composition effect, return to the stolen data market. Consider the price legitimate sellers can expect to receive from a given data package and how that price changes with r :

$$p^* = \frac{S^2V}{(S+V)^2} = \frac{r^2V}{(1+r)^2}$$

$$\frac{\partial p}{\partial r} = \frac{2rV(1+r) - 2r^2V(1+r)}{(1+r)^4} = \frac{2rV}{(1+r)^3}$$

For $r > 0$, the r values of interest, price is increasing in r . Price increases in r because q , the portion of legitimate sellers, is also increasing in r (figure 3):

$$q = \frac{S}{S+V} = \frac{r}{1+r}$$

$$\frac{\partial q}{\partial r} = \frac{1}{(1+r)^2}$$

As r increases, the portion of sellers that are legitimate increases. This increases demand for the good, raising prices and incentivizing more breaches of data-rich targets. Effectively, the regulation alleviates the stolen data market's lemons problem enough to bring more large data breaches into the market.

Summarizing the above, the model makes the following predictions about each outcome of

interest.

Total number of breaches will decrease if the number of data-rich targets is small enough. The probability the most data-rich targets are breached increases after the policy because of the downstream effects in the stolen data market. In that market, privacy regulations reduce the equilibrium number of scammers, increasing the price of the data packages taken from these larger targets. However, breach probability falls or remains constant for all other targets.

The average breach size effect depends on two factors. First, the increase in breach probability for the most data-rich and decrease in breach probability for all other targets changes the composition of stolen data packages and increases expected breach size. But the reduction in data collection by all targets decreases expected breach size. The overall effect will depend on whether the composition or data reduction effect dominates.

The total amount of data effect will similarly depend on whether the composition or data reduction effect dominates. If the composition effect dominates and the most data-rich targets collect more data post-regulation than the targets whose post-regulation breach probability decreases did prior to the regulation, the total amount of data in the market will increase.

4 Stolen Data Market Observations

Data for this study come primarily from SpyCloud, a private cybersecurity company specializing in identity threat protection. They have constructed a catalog of data breaches gathered from online stolen data marketplaces. Each observation is of a data package posted in the market and contains information on where the data originated (either the country or specific organization), what types of data were stolen, and the total number of records included in the breach. The data packages were available online between 2015 and 2023, though the breaches they originate from may have occurred as early as 2002. To the best of my knowledge, this is the first time such a dataset has been used to quantitatively study the effect of any policy change on the stolen data market. The details of the data allow me to go deeper than the aggregate and summary statistics previous research has depended on to see what is actually being traded. Unfortunately, I do not observe prices for all but a handful of data packages, restricting this paper to measuring only quantity effects.

Panel A of table 2 displays summary statistics for the three outcomes of interest in the study at the data package level: the total number of records in a data package; the fraction of the data in a data package that is personally identifiable information (PII); and the number of data types in each package. Examples of observed data types are email addresses, credit card information, or whether the identified person owns a cat. PII has multiple legal definitions, but is broadly any information that can identify an individual and may not be publicly known. A more in-depth discussion of the definition of PII is in section 1 of the online appendix.

Data packages vary greatly in terms of size, measured by the number of records. The largest contains over 700 million data points, while the smallest contains only one. Similarly, they range from having only one type of data to 55. Where they are more alike is in the fraction of records

Table 2: Data Package Summary Statistics

Panel A: Data Package Summary Statistics

	Number of Records	PII Fraction	Number of Data Types
Observations	4,394	4,394	4,394
Mean	3,544,186	0.693	6.220
Std. Dev.	28,996,342	0.191	5.208
Min.	1	0	1
25%	5,164	0.500	2
50%	46,748	0.667	4
75%	288,555	0.858	9
Max.	716,409,393	1	55

Panel B: Group Counts

	Pre-GDPR N=1,621		Post-GDPR N=2,773	
	Non-EU	EU	Non-EU	EU
N	1,175	446	2,292	481

Note: PII fraction is the fraction of records in a data package that are considered PII. A discussion of what constitutes PII is in the appendix.

in the breach that are personally identifiable information. The 25th percentile breach is 50 percent PII, and the 75th percentile breach has 85 percent PII. Emails, considered PII under the GDPR, are the most common type of data in these breaches, closely followed by passwords (figure 4).

Because the GDPR applies to any entity collecting data on EU residents, not just those located in the EU, identifying treatment and control groups is difficult. For each data package used in the analysis, I observe either the country from which the data originate or the name of the organization that was breached, and both for a subset of the observations. When I observe only the originating country, I assign the breach to that country. This makes treatment categorization simple: if the data originates in the EU, the package is treated. In the cases where I only observe the organization from which the data were stolen, I use one of two processes. First, I attempt to determine where the organization is headquartered. If they are an EU-based organization, the package is treated. If they are not, I search their website (where available) to see whether they indicate having a European presence. Those that do are categorized as treated. In addition to the manual search, I use a method similar to Demirer et al. (2024), who use firms' data server locations to categorize them into treatment and control units, to determine country. I cannot observe data center locations, so instead I use the locations of the servers on which they host their websites, as that location is endogenous to the location of an organization's users and customers.⁹

⁹Although it has largely been abstracted away, the internet is fundamentally a physical network. Data flows through fiber optic cables that span across oceans and continents to deliver content to users. This means the further a user is physically located from the server hosting the content, the longer it takes content to be delivered.




FIGURE 4
Fraction of Data Packages Containing Each Data Type

Note: The password category includes both individual passwords themselves, and information related to passwords, such as the salt used to help obscure them. Financial information includes bank, credit card, and loan data. Treated refers to all data packages originating in the EU, untreated to those originating outside the EU.

There are two pieces of the internet’s architecture key to connecting users with websites that allow me to observe where sites are physically located: DNS and GeoDNS. A Domain Name System (DNS) is essentially a phonebook for the internet. When a user types a domain name (e.g., www.fangraphs.com) into their web browser, it sends a query to the DNS, which then finds the IP address of the server hosting that website and connects it to the user. A GeoDNS does the same while taking into account the location of the user sending the initial query. For websites hosted in multiple locations, it will respond with the IP address of the server hosting the requested website that is closest to the user. As an example, suppose a website is hosted on one server in San Francisco, California, and another in Berlin, Germany. A user in Los Angeles will be connected to the San Francisco server, and a user in Frankfurt will be connected to the Berlin server.

To find where an organization hosts their website, I use the GeoNet API tool from Shodan, an internet devices research company.¹⁰ The GeoNet API allows me to send GeoDNS queries from six locations around the world to any website and record the IP addresses that respond to each request.¹¹ I conduct these queries for each data package in my sample with an associated organizational website. After collecting the IP addresses of the responding servers, I use Shodan’s IP address lookup tool to find the physical location of each one. Under this method, I categorize a breach as having come from a regulated entity if the organization hosts their website on at least one server in the EU. An organization that hosts their website both in and outside the EU will also be considered regulated. Data packages that have not been manually assigned to a specific country are assigned to the country in which their originating organization hosts a majority of their servers.

The difference in time may only be fractions of a second, but that can still have a noticeable effect on outcomes organizations care about. Previous research has found that a 0.1 second improvement in website load time can increase spending on retail sites by almost 10% (Deloitte, 2020). For streaming and gaming sites, decreasing lag time improves user experience and can be used as a differentiating factor. There is a clear incentive for organizations to host their website on servers that are physically near their users to minimize load time.

¹⁰<https://geonet.shodan.io>

¹¹Requests are sent from servers in the United States, England, the Netherlands, Germany, India, and Singapore.

Panel B of table 2 breaks down the number of data packages that fall into each category before and after the GDPR.

Data packages are only included in the final dataset if I can determine whether the originating organization is subject to the GDPR (or was in the cases where the organization is no longer active), and when the data are available online (pre- or post-GDPR). This sample represents only data packages that are posted online, not necessarily all data that is collected or stolen. It is possible that some stolen data packages are not traded, in which case I cannot observe them.¹² There are many reasons why a package may be unobservable. The hacker may decide they can profit more from using the data themselves than from publishing it. Or the hacker may have full access to the organization’s data, but decide only a subset is worth taking and selling. In the case of ransomware, the victim organization may decide to pay the ransom to prevent their data from being leaked.


Using the country assigned to each package and the date it was available online, I aggregate the individual packages into a country-quarter level panel. The panel spans January 2017 to November 2023. Each country can be thought of as a market in the theoretical model in Section 3. The value and cost of hacking organizations in regulated countries will be affected by the GDPR, and will remain unchanged in unregulated countries. Choosing January 2017 to November 2023 as the study’s time frame means that data packages available online prior to 2017 are not included in the panel, even though I can observe full information on them. I make the choice to exclude pre-2017 periods because SpyCloud was started in 2016. This padding removes any bias that may occur if the packages collected early in their operation are fundamentally different from the ones discovered later. For consistency in the sample I also exclude these observations from the primary data package level analysis. To control for population in the analysis, I add annual population data from the World Bank to the panel.

For robustness checks, I construct additional panels excluding any period after March 2020—to remove any bias introduced by the COVID-19 pandemic, and any data package originating from a multinational organization. The latter removes any bias that may arise due to partially treated organizations.


Not every country experiences a breach in every period. For those observations, I assign a value of zero to the two outcome variables: number of data breaches and number of records. As shown in figures 5 and 6, this creates mass points at zero for both variables. I discuss the implications of this for my estimation strategy in section 5.

Roughly 27 percent of all country-quarters have a positive number of breaches (panel A of table 3), and among the positive observations there are an average of six breaches and 21.6 million records stolen (panel B of table 3). There is a large variation in both outcomes with as many as 245 breaches occurring and over one billion records being available in a quarter. Figures 7a and 7b show how the number of data breaches and number of records trended over time. There is a clear decline in the number of data breaches immediately after the GDPR went into effect, but no


¹²If only part of a data package is traded, I only observe what is traded, not everything that was stolen.




(A) Logged Using All Observations



(B) Logged Using Only Positive Observations




(C) In Levels Using All Observations




(D) In Levels Using Only Positive Observations

FIGURE 5
Distribution of the Aggregate Number of Data Breaches Per Country by Quarter


obvious and persistent change in the number of records becoming available in each quarter. The fraction of total breaches originating in the EU also falls following the GDPR (figure 8).




(A) Distribution of Log(Number of Records) Using All Observations



(B) Distribution of Log(Number of Records) Using Only Positive Observations




(C) Distribution of Number of Records Using All Observations




(D) Distribution of Number of Records Using Only Positive Observations

FIGURE 6
Distribution of the Aggregate Number of Records Per Country by Quarter



(A) Number of Data Breaches



(B) Number of Records Online

FIGURE 7
Number of Breaches and Records Online Time Series

Table 3: Panel Data Summary Statistics

Panel A: All Periods

	Number of Breaches	Number of Records (M)	> 0 Breaches
Observations	2,716	2,716	2,716
Mean	1.62	5.73	0.27
Std. Dev.	9.52	49.73	0.44
Min.	0	0.00	0
25%	0	0.00	0
50%	0	0.00	0
75%	1	0.00	1
Max.	245	1,009.74	1

Panel B: Non-Zero Periods Only

	Number of Breaches	Number of Records (M)
Observations	721	721
Mean	6.09	21.60
Std. Dev.	17.74	94.78
Min.	1	0.00
25%	1	0.04
50%	2	0.26
75%	4	1.90
Max.	245	1,009.74




FIGURE 8
Fraction of Breaches That Originate in EU

5 Empirical Strategy

I estimate the effects of the GDPR on aggregate quantities in the stolen data market, and the contents of the individual data packages traded. The model I propose in section 3 predicts that there will be a decrease in the total number of data packages originating in the EU. The change in the total amount of data available and average size of a data breach is ambiguous.

Treatment status in all cases is determined by where the data was originally collected and the date the data package was available online. A data package is in the treatment group if it originated in the EU or was stolen from an organization subject to the GDPR and became available in June 2018 or later. This definition includes multinational organizations, such as large social media organizations, as treated if they have any users in the EU. As discussed in Demirer et al. (2024), this complicates identification because these organizations may respond differently to the GDPR. The data they hold is partially treated since they likely hold information on individuals inside and outside the EU. Robustness checks excluding data packages from multinationals are in section 4.2 of the online appendix, and their findings are discussed in sections 6.1 and 6.2.

5.1 Number of Data Packages and Quantity of Data Available

The two aggregate effects of interest are the change in the number of data packages and total amount of data available. They are estimated using the country-quarter panel described in section 4. Each observation of country i is the aggregate of the individual data packages originating from that country in time period t . Most countries do not have a positive number of breaches in each period, creating a mass point at zero (figures 5 and 6). Privacy regulations are expected to affect the extensive margin because they change the relative value of breaching organizations in regulated countries, making them less likely to have a positive number of breaches in a given period. To measure the extensive margin effect, I estimate the linear probability model:

$$Positive_{it} = \gamma_i + \tau_t + \delta D_i \times Post-GDPR_t + \varepsilon_{it}.$$

where γ_i and τ_t are country and quarter fixed effects. D_i equals one if the country is in the EU, and $Post-GDPR_t$ equals one if the period is after the second quarter of 2018. The dependent variable, $Positive_{it}$, is an indicator for whether country i has at least one breach in period t .

To measure the impact of the GDPR on the number of breaches and total number of records available, I use a Poisson model:

$$Y_{it} = \exp(\gamma_i + \tau_t + \delta^{Agg} D_i \times Post-GDPR_t - \log(population_{it})) \varepsilon_{it} \quad (8)$$

where γ_i , τ_t , D_i , and $Post-GDPR_t$ are all defined as in the extensive model. To explicitly obtain the percentage change in the outcome, δ^{Agg} must be transformed:

$$\delta^{Agg\%} = \exp(\delta^{Agg}) - 1. \quad (9)$$

Standard errors are clustered at the country level. The offset, $\log(population)$, is used to account for difference in sizes among the countries.

I use a Poisson model rather than a log-like transformation because of the mass points at zero. In order to use log-like transformations on the outcomes, it would be necessary to either add a constant to each observation or use a transformation that is defined at zero, such as the inverse hyperbolic sine, to include the full sample in the estimation. Mullahy and Norton (2024) show that log-like transformations significantly change the estimated marginal effects when zero mass points are present. Further, J. Chen and Roth (2023) find that, in the presence of zero mass points, if the treatment has extensive margin effects, the estimated average treatment effect is sensitive to the units of the outcome variable, making the interpretation of the estimates difficult. The model I present in section 3 makes clear that privacy regulations should affect the extensive margin as it changes the probability organizations in regulated countries are breached, changing the likelihood those countries experience a positive number of breaches.

To test whether the effect changes over time, I break the $Post-GDPR_t$ term into short- and long-run effects, estimating:

$$Y_{it} = \exp\left(\gamma_i + \tau_t + \delta_{SR}^{Agg} Short-Run_t \times D_i + \delta_{LR}^{Agg} Long-Run_t \times D_i - \log(population_{it})\right) \varepsilon_{it} \quad (10)$$

where $Short-Run_t$ equals one when $t \in \{\text{June 2018} - \text{May 2019}\}$ and $Long-Run_t$ equals one for all periods after May 2019.¹³ This allows for changes in the behavior of both those collecting and stealing data that lags behind the policy implementation. Organizations may increase their compliance level, and hackers may reevaluate their strategies after observing organizational behavior changes.

The identifying assumptions underlying these models are conditional no anticipation, and that the growth rate between periods the treated group would have realized in the absence of treatment is the same as that experienced by the control group, i.e., there are parallel trends in the ratio of outcomes between periods (Wooldridge, 2023):

$$\frac{E[Y^{Post}(0)|D=1]}{E[Y^{Pre}(0)|D=1]} = \frac{E[Y^{Post}(0)|D=0]}{E[Y^{Pre}(0)|D=0]}.$$

To test this assumption, I estimate an event study model:

$$Y_{it} = \exp\left(\gamma_i + \tau_t + \sum_{t \neq -1} \delta_{it}^{Agg} D_i \times Post-GDPR_t - \log(population)\right) \varepsilon_{it} \quad (11)$$

where all notation is defined as before and standard errors are once again clustered at the country level.

Under the model in section 3, the increase in cost and decrease in value of breaching regulated

¹³The short- and long-run definitions follow Demirer et al. (2024).

organizations caused by the GDPR should cause the number of data breaches originating in regulated countries to decrease. All else equal, the number of records should decrease as well, but changes in which targets are hacked and which data packages are subsequently sold may blunt this effect. If there is a sufficiently large shift towards data-rich targets post-GDPR, the expected number of records may even increase.

5.2 Data Package Content Effects

Effects on individual data packages are estimated using the linear model:

$$y_j = \gamma_j + \tau_t + \delta^{DP} D_j \times Post-GDPR_t + \varepsilon_j \quad (12)$$

where D_j equals one if the package originated from a regulated organization, and $Post-GDPR_t$ indicates whether the data package was available June 1, 2018 or later. I use June 1, 2018 as the beginning treatment date, rather than the day the GDPR was enforced, to allow for a lag between when data became available online and when it was stolen.

The three outcomes of interest are the log of the total number of records in the package, the amount of personally identifiable information (PII), and the number of unique types of data in the package. The parameter of interest is δ^{DP} .

I once again break the $Post-GDPR$ term into short- and long-run effects and estimate:

$$y_j = \gamma_j + \tau_t + \delta_{SR}^{DP} D_j \times Short-Run_t + \delta_{LR}^{DP} D_j \times Long-Run_t + \varepsilon_j \quad (13)$$

where $Short-Run_t$ and $Long-Run_t$ are defined as they were in the aggregate effects section.

The expected effects on individual data packages are ambiguous under the model in section 3. If the GDPR increases the portion of breaches originating from data-rich organizations enough, then the expected amount of data in a data breach increases. Given the restrictions on collecting PII, the fraction of all records that are PII might be expected to decrease, but that will depend on the effects of the regulation on non-PII data collection as well.

6 Results

The main results are presented in sections 6.1 and 6.2. Robustness checks and alternative model specifications are discussed in section 6.3.

6.1 Number of Data Packages and Quantity of Data Available

On the extensive margin, I find the GDPR is associated with a roughly 20 percent decline in the probability of finding a data breach that originates from a regulated country online (table 4). This effect is larger in the long-run than in the short-run (-21.1 percent versus -15.8 percent, respectively).

Table 4: Extensive Margin Effects

	Dependent Variable: Positive Number of Breaches	
	(1)	(2)
Post $\times$ Treatment	-0.201*** (0.039)	
SR $\times$ Treatment		-0.158*** (0.049)
LR $\times$ Treatment		-0.211*** (0.039)
Observations	2,716	2,716
R^2	0.469	0.469
Period Fixed Effects	Y	Y
Country Fixed Effects	Y	Y

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Note: Standard errors are clustered at the country level.

The aggregate treatment effects on the number of breaches and total amount of data posted from a country are presented in table 5. I find that the number of data breaches fell 54.8 percent and 60.9 percent in the short and long run, respectively, and 60.4 percent overall. This result is consistent with the predicted effects of both a decrease in the amount of data collected and an increase in security investment by regulated organizations on the market. Fewer organizations are worth hacking, so there is a decrease in the number of data breaches. The same logic applies to my extensive margin findings.

Despite the large decrease in the number of breaches, I find no statistically significant change in the number of records in the market. Interestingly, the short- and long-run effect estimates indicate there may have been an initial decrease in the aggregate number of records available, followed by an increase in the long run. Intuitively, this could be the result of hackers making adjustments to their strategies in the long run. A shift in strategy towards higher-valued targets would result in larger data packages becoming a larger share of observed postings, which is mechanically necessary for these results to be true.

Event study plots to provide evidence that the parallel trends assumption holds are in figure 9. For number of data breaches, the coefficient estimates for each period prior to the GDPR have zero in the 95 percent confidence interval, while post-GDPR there is a clear decrease in the number of data breaches (figure 9a). Each period of the event study shows no statistically significant effect on the number of records traded (figure 9b).

6.2 Individual Data Package Content Effects

At the individual breach level, I find that data packages originating in regulated organizations increased in size by 62.9 percent, as measured by number of records they contain (column three

Table 5: Aggregate Effects

	Number of Breaches		Number of Records	
	(1)	(2)	(3)	(4)
Post $\times$ Treatment ($\hat{\delta}^{Agg}$)	-0.927*** (0.266)		0.342 (0.429)	
SR $\times$ Treatment ($\hat{\delta}_{SR}^{Agg}$)		-0.794*** (0.297)		-0.221 (0.590)
LR $\times$ Treatment ($\hat{\delta}_{LR}^{Agg}$)		-0.939*** (0.284)		0.407 (0.429)
$\hat{\delta}^{Agg\%}$	-0.604 (0.105)		0.407 (0.604)	
$\hat{\delta}_{SR}^{Agg\%}$		-0.548 (0.134)		-0.198 (0.473)
$\hat{\delta}_{LR}^{Agg\%}$		-0.609 (0.111)		0.502 (0.645)
Period Fixed Effects	Y	Y	Y	Y
Country Fixed Effects	Y	Y	Y	Y
Observations	2,716	2,716	2,716	2,716
Pseudo R^2	0.792	0.793	0.847	0.847

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Note: Standard errors are clustered at the country level. $\hat{\delta}^{Agg\%}$ terms are calculated by transforming the $\hat{\delta}^{Agg}$ terms. $\hat{\delta}^{Agg\%} = \exp(\hat{\delta}) - 1$. The standard errors for these terms are calculated using the Delta Method: $SE(\hat{\delta}^{Agg\%}) = \exp(\hat{\delta}^{Agg})SE(\hat{\delta}^{Agg})$.

of table 6). This effect is driven by long-run changes, with there being no statistically significant short-run effect (column four of table 6).

An increase in the size of the data packages is counterintuitive on its face. If data privacy legislation successfully reduces data collection, which it appears to do, then it seems natural that there would be a corresponding reduction in the number of records included in the packages. Less data collected means there is less data to steal. However, if the GDPR brings more high-value breaches into the market, then the expected value of the remaining breaches increases even after accounting for the change in value caused by the GDPR. These breaches would presumably contain larger amounts of data, increasing the expected number of records in any given breach. Figure 10 shows that the distribution of the number of records in a breach shifted right after the GDPR for both the treated and untreated groups. The main difference is in the tail end of distributions. The untreated group appears to have relatively more weight at the lower end of the distribution, while the treated group has more weight at the very high end.

Looking specifically at the amount of PII in a breach, I find that the number of records that constitute PII increased by 60.6 percent in the long run (column 8 of table 6). Given that most of the data in the packages qualifies as PII (table 2), this is expected with the increase in the overall




FIGURE 9
Aggregate Effect Event Studies

Note: The figures present estimates of the δ_{it}^{Agg} coefficients in equation 11 converted to percentage changes using $\exp(\delta_{it}^{Agg}) - 1$. The bars are the 95 percent confidence intervals with standard errors clustered at the country level. Period $t = -1$, the first quarter of 2018, is normalized to be zero.

number of records per package.

These are the only statistically and economically significant changes at the data package level. I find no change in the fraction of records that are PII (columns 1-4 of table 7) or number of unique types of data in the packages (columns 5-8 of table 7). One potential explanation for this is that only certain types of data have value in the stolen data market. If the data regulated organizations ceased collecting after the GDPR is not considered valuable in this other market, it is unlikely that there would be an effect on data package contents beyond their size.




FIGURE 10
Number of Records Per Data Package Density

Note: The distributions shown are of the number of records in each data package. This differs from figure 5, which is the distribution of the aggregate number of records observed for each country in all periods.

Table 6: Data Package Effects: Number of Records

	Log(Total Number of Records)				Log(Number of PII Records)			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Post $\times$ Treatment ($\hat{\delta}^{DP}$)	0.936** (0.412)		0.488* (0.251)		0.933** (0.422)		0.441* (0.247)	
SR $\times$ Treatment ($\hat{\delta}_{SR}^{DP}$)		0.691 (0.491)		0.196 (0.333)		0.677 (0.542)		0.134 (0.384)
LR $\times$ Treatment ($\hat{\delta}_{LR}^{DP}$)		0.963** (0.414)		0.519** (0.258)		0.960** (0.422)		0.474* (0.251)
Multinational			1.390*** (0.252)	1.392*** (0.252)			1.525*** (0.268)	1.527*** (0.269)
$\hat{\delta}^{DP\%}$	1.550 (1.051)		0.629 (0.408)		1.542 (1.073)		0.555 (0.384)	
$\hat{\delta}_{SR}^{DP\%}$		0.996 (0.981)		0.216 (0.405)		0.969 (1.068)		0.143 (0.439)
$\hat{\delta}_{LR}^{DP\%}$		1.618 (1.085)		0.680 (0.433)		1.613 (1.103)		0.606 (0.404)
Period Fixed Effects	Y	Y	Y	Y	Y	Y	Y	Y
Country Fixed Effects	Y	Y	Y	Y	Y	Y	Y	Y
Observations	4,394	4,394	4,394	4,394	4,394	4,394	4,394	4,394
R^2	0.268	0.268	0.276	0.276	0.270	0.270	0.279	0.279

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Note: Standard errors are clustered at the country level. $\hat{\delta}^{DP\%}$ terms are calculated by transforming the $\hat{\delta}^{DP}$ terms. $\hat{\delta}^{DP\%} = \exp(\hat{\delta}) - 1$. The standard errors for these terms are calculated using the Delta Method: $SE(\hat{\delta}^{DP\%}) = \exp(\hat{\delta}^{DP})SE(\hat{\delta}^{DP})$. Multinational is an indicator variable equal to one if the data package originated from an organization active in multiple countries.

Table 7: Data Package Effects: Content

	PII Fraction				Number of Unique Data Types			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Post $\times$ Treatment ($\hat{\delta}^{DP}$)	0.002 (0.007)		-0.005 (0.011)		0.420* (0.225)		0.392 (0.245)	
SR $\times$ Treatment ($\hat{\delta}_{SR}^{DP}$)		-0.013 (0.028)		-0.020 (0.023)		0.263 (0.586)		0.232 (0.585)
LR $\times$ Treatment ($\hat{\delta}_{LR}^{DP}$)		0.003 (0.008)		-0.004 (0.013)		0.437* (0.260)		0.409 (0.279)
Multinational			0.021 (0.020)	0.022 (0.020)			0.086 (0.231)	0.087 (0.231)
Period Fixed Effects	Y	Y	Y	Y	Y	Y	Y	Y
Country Fixed Effects	Y	Y	Y	Y	Y	Y	Y	Y
Observations	4,394	4,394	4,394	4,394	4,394	4,394	4,394	4,394
R^2	0.423	0.423	0.424	0.424	0.280	0.280	0.280	0.280

* $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$

Note: Standard errors are clustered at the country level. PII fraction is the number of records in a data package considered PII divided by the total number of records in the data package.

6.3 Robustness

To check the robustness of my results, I re-estimated the aggregate effects using a number of alternative samples and model specifications. Splitting the sample into small and large countries, defined by whether the population is above or below the median in 2018, I find that the point estimate is slightly more negative for small countries than for large countries, but the directional effects are the same. The only notable change is that the short-run effect in large countries is statistically insignificant.¹⁴

The effects on the number of data packages and total quantity of data available also hold under different model and panel construction choices. None of removing the population offset, weighting the observations by population, or changing the outcome variables to per capita variables changes the direction or significance of the results.¹⁵ Removing all observations from Brazil and China, which each implemented privacy regulations of their own near the end of the study, has no effect. Nor does removing all multinational breaches from the aggregate numbers or the post-COVID-19 periods.

Using log-like transformations, $\text{Log}(Y + 1)$ and the inverse hyperbolic sine, as alternatives to the Poisson model changes the magnitude of the effect on both aggregate outcomes and shows there is a significant decrease in the number of records available. However, these outcome changes are all consistent with the bias J. Chen and Roth, 2023 and Mullahy and Norton, 2024 show log-like transformations introduce when used on data with a mass point at zero.

The data package content effects are also robust to changes in how the sample is constructed.

A more in-depth discussion of these robustness checks can be found in the online appendix.

7 Discussion and Conclusion

As organizations continue to collect large amounts of data, the risk of that data being stolen and sold will be ever-present. In this paper I have shown that data protection regulations can have a significant effect on the illicit market for data.

I estimate that the GDPR is associated with a 60 percent reduction in the number of data breaches originating in EU countries available in stolen data markets. There is, however, no accompanying reduction in the number of individual records in these markets, as the size of data packages increased over 60 percent as well. I find no other changes in the contents of the data packages.

These results are consistent with what the model in section 3 predicts will happen after a privacy regulation goes into effect. Although a privacy regulation reduces breach probability for some legal data collectors, it will increase breach probability for the extremely data-rich legal data collectors. This is because the regulation alleviates the lemons problem in the stolen data market, incentivizing

¹⁴I define a country as small or large based on whether their population is below or above the median country population in 2018.

¹⁵It should be noted that converting the outcomes to per capita terms changes them from discrete to continuous variables, making a Poisson model inappropriate to use.

more hacks against larger targets. The result is that even though the expected number of breaches falls, the expected total amount of data and average breach size may increase.

This paper partially fills the gap in the literature on privacy regulation, and the GDPR in particular, regarding potential benefits of these regulations. It is the first to study the effects of privacy regulation on the stolen data market and show a causal impact.

There are many paths forward for future research on this topic. Qualitative and quantitative work on the abilities and behaviors of hackers could provide insights into whether the effects I observe empirically are driven strictly by the changes in hacker incentives I propose, or if there are other factors, such as changes in hacker skill, at play.

This paper is missing a key component of the market: prices. Although there are many hurdles to collecting quality price data in these markets, doing so would open the door to a more complete analysis of their inner workings and the value hackers place on certain types of information.

Implicit in my model is the assumption that hacker skill remains constant. Human capital accumulation on the part of hackers may be an equilibrium response to the GDPR. If hackers were to become more productive, the cost of hacking would decrease, resulting in more breaches, but the expected value may decrease as relatively low-value targets become viable marks now that they are cheaper to hack. That I find a decrease in the number of breaches and an increase in their quality (in terms of size) suggests this is not a concern. However, I do not observe any direct measure of hacker ability and therefore cannot fully rule out the possibility that it has changed.

Finally, while there have been a number of studies on the effects of the GDPR on specific firm outcomes, changes in data collection, and now cybercrime, there is still no overarching study of its overall welfare effects or how individual components of the policy influence outcomes of interest. Reducing the number of data breaches is likely beneficial to all not looking to buy or sell stolen data. But the extent of that benefit may be limited given that the number of records in stolen data markets did not also fall. With at least the same amount of data available if not more, individuals may be no better off than they were before. To test this, one would need to calculate how the probability of a person's data being online has changed, or at least count the number of unique individuals with data in each breach, which I am unable to do with my data.

On the cyber criminal side of the problem, the GDPR may have made them better off in some cases. By incentivizing a shift in effort towards higher-value targets and potentially alleviating the stolen data market's asymmetric information problem, the GDPR may have provided some benefit to internet miscreants.

A thorough welfare analysis will also need to account for the externalities of data collection (Choi, Jeon, and Kim, 2019; Acemoglu et al., 2022), and the larger costs of cyberattacks, such as ransomware attacks on hospitals (Neprash, McGlave, and Nikpay, 2026).

For policymakers, a key question is whether subsequent privacy regulations will have effects similar to the GDPR, or if the reduction in breaches against the EU is really the result of cybercrime being displaced towards non-regulated targets. A spillover could be expected if hackers are effort-constrained and rank all potential targets, breaching only the top choices until they exhaust their

effort endowment. Under this model, unregulated organizations become relatively more valuable targets, increasing the probability they are breached. The reduction in data breaches I observe would be a combination of fewer breaches against regulated organizations, and an increase in breaches against non-regulated ones. While I am not able to directly test for spillovers, an increase in the number of breaches against non-regulated targets would appear in the time series plots around the GDPR's passage. As can be seen in figure [7a](#), no such increase is apparent. This suggests there is no substantial displacement effect. With more countries considering and adopting data privacy regulations, research on this subject would have high returns in the debate over how to design future policy.

Appendix

A Model Derivations

A.1 The Stolen Data Market

A.1.1 Demand

Demand in the stolen data market equals the probability a buyer finds a data package worth purchasing at a given price:

$$\Pr(\kappa\mu > p)$$

With the simplifying assumption that $\kappa \sim \text{Uniform}[0, 1]$, and the behavioral assumption that buyers will only purchase a data package if the expected utility is positive:

$$Q^D(p, \mu) = \Pr(\kappa\mu > p) = \Pr(\kappa > \frac{p}{\mu}) = 1 - \frac{p}{\mu}$$

A.1.2 Legitimate Supply

Legitimate supply in the market is provided by heterogeneous hackers with utility

$$U_h = p - \gamma S$$

As with demand, use the simplifying assumption that $\gamma \sim \text{Uniform}[0, 1]$ and the behavioral assumption that sellers only offer packages they expect to at least break even on, with a unit mass of legitimate hackers, total legitimate supply is:

$$Q_\ell^S(p, S) = \frac{p}{S}$$

A.1.3 Scammer Optimization

The profit maximizing representative scammer solves:

$$\max_q (1 - q) \left(1 - \frac{p}{qV} \right) p$$

Taking the first order condition and solving for q ,

$$\begin{aligned}
-1 \left(1 - \frac{p}{qV}\right) p + (1 - q) p \left(\frac{p}{q^2 V}\right) &= 0 \\
\rightarrow q^* &= \sqrt{\frac{p}{V}}
\end{aligned}$$

The equilibrium price and q are derived as in the main body of the paper:

$$p^* = \frac{S^2 V}{(S + V)^2} \quad q^* = \frac{S}{S + V}$$

A.2 Legal Data Collectors

Legal data collector's breach probability is equal to the probability a legitimate hacker finds it profitable to target them. Using the supply curve and equilibrium price, this is:

$$\pi(S, V) = \frac{SV}{(S + V)^2}$$

Their loss function is

$$L(S, V) = \frac{\Lambda V^\psi}{\delta S}, \quad 0 < \psi < 2$$

Substituting $S = rV$ and simplifying yields a simple function for expected losses:

$$\mathbb{E}[Loss] = \pi(r, V)L(r, V) = \frac{rV^2}{(rV + V)^2} \frac{\Lambda V^\psi}{\delta rV} = \frac{\Lambda}{\delta(1 + r)^2} V^{\psi-1}$$

With two types of data, the legal data collector optimization problem is

$$\max_{r, d_1, d_2} A (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\phi/\rho} - \frac{\Lambda}{\delta(1 + r)^2} (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\frac{\psi-1}{\rho}} - \sum_{j=1}^2 \omega_j d_j - \omega_S r (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{1/\rho}$$

Taking the first order conditions with respect to the d terms:

$$\begin{aligned}
d_1 : \left[\frac{\phi}{\rho} A V^{\phi-\rho} - \frac{\Lambda}{\delta(1 + r)^2} \frac{\psi-1}{\rho} V^{\psi-1-\rho} - \omega_S r \frac{1}{\rho} V^{1-\rho} \right] \alpha_1 \rho d_1^{\rho-1} &= \omega_1 \\
d_2 : \left[\frac{\phi}{\rho} A V^{\phi-\rho} - \frac{\Lambda}{\delta(1 + r)^2} \frac{\psi-1}{\rho} V^{\psi-1-\rho} - \omega_S r \frac{1}{\rho} V^{1-\rho} \right] \alpha_2 \rho d_2^{\rho-1} &= \omega_2
\end{aligned}$$

Taking the ratio of the two,

$$\begin{aligned}
\frac{\alpha_1 d_1^{\rho-1}}{\alpha_2 d_2^{\rho-1}} &= \frac{\omega_1}{\omega_2} \\
\rightarrow d_1 &= \left[\frac{\omega_1}{\omega_2} \frac{\alpha_2}{\alpha_1} \right]^{\frac{1}{\rho-1}} d_2
\end{aligned}$$

Plugging this into the value production function yields the conditional factor demand for d_1 :

$$\begin{aligned}
V &= (\alpha_1 d_1^{*\rho} + \alpha_2 d_2^{*\rho})^{1/\rho} \\
V^\rho &= \underbrace{\left[\alpha_1 \left[\frac{\omega_1}{\omega_2} \frac{\alpha_2}{\alpha_1} \right]^{\frac{\rho}{\rho-1}} + \alpha_2 \right]}_{\equiv B_2} d_2^\rho \\
\rightarrow d_2(V) &= \frac{V}{B_2^{1/\rho}}
\end{aligned}$$

By symmetry,

$$d_1(V) = \frac{V}{B_1^{1/\rho}}$$

where

$$B_1 = \left[\alpha_2 \left[\frac{\omega_2}{\omega_1} \frac{\alpha_1}{\alpha_2} \right]^{\frac{\rho}{\rho-1}} + \alpha_1 \right]$$

Taking the first order condition with respect to r will then give r as a function of V :

$$\begin{aligned}
\frac{2\Lambda}{\delta(1+r)^3} V^{\psi-1} &= \omega_S V \\
\left[\frac{2\Lambda}{\delta\omega_S} \right] V^{\psi-2} &= (1+r)^3 \\
r(V) &= \left[\frac{2\Lambda}{\delta\omega_S} \right]^{\frac{1}{3}} V^{\frac{\psi-2}{3}} - 1
\end{aligned}$$

To find the optimal V , consider the total cost of data collection:

$$\begin{aligned}
TC(V) &= \omega_1 d_1(V) + \omega_2 d_2(V) + \omega_S r(V) V + \frac{\Lambda}{\delta(1+r(V))^2} V^{\psi-1} \\
&= \omega_1 d_1(V) + \omega_2 d_2(V) + \omega_S \left[\frac{2\Lambda}{\delta\omega_S} \right]^{\frac{1}{3}} V^{\frac{\psi+1}{3}} - \omega_S V + \frac{\Lambda}{\delta \left[\frac{2\Lambda}{\delta\omega_S} \right]^{2/3}} V^{\frac{\psi+1}{3}}
\end{aligned}$$

The marginal cost of V is then

$$MC(V) = \frac{\omega_1}{B_1^{1/\rho}} + \frac{\omega_2}{B_2^{1/\rho}} + \left(\omega_S \left[\frac{2\Lambda}{\delta\omega_S} \right]^{1/3} + \frac{\Lambda}{\delta \left[\frac{2\Lambda}{\delta\omega_S} \right]^{2/3}} \right) \frac{\psi+1}{3} V^{\frac{\psi-2}{3}} - \omega_S \quad (14)$$

The marginal benefit of an additional unit of V is

$$\frac{\partial}{\partial V} AV^\phi = MB(V) = \phi AV^{\phi-1} \quad (15)$$

The optimal V is found by setting the marginal cost (equation 14) equal to the marginal benefit (equation 15), which does not yield a closed form solution without additional assumptions over ϕ .

After the GDPR, the firm's optimization problem includes a constraint term:

$$\begin{aligned} \max_{r, d_1, d_2} A (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\phi/\rho} - \frac{\Lambda}{\delta(1+r)^2} (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{\frac{\psi-1}{\rho}} - \sum_{j=1}^2 \omega_j d_j - \omega_S r (\alpha_1 d_1^\rho + \alpha_2 d_2^\rho)^{1/\rho} \\ + \lambda [r - \underline{r}] \end{aligned}$$

Taking the ratio of the first order conditions with respect to d_1 and d_2 results in the same equation and conditional factor demands as the unconstrained problem. Taking the first order condition with respect to r ,

$$\frac{2\Lambda}{\delta(1+r)^3} V^{\psi-1} - \omega_S V + \lambda = 0$$

When the constraint is not binding, $\lambda = 0$ and the solution is the same as the unconstrained solution. When the constraint does bind, $r = \underline{r}$. This changes the marginal benefit and marginal cost comparison. For constrained targets, the total cost of V is

$$TC(V) = \omega_1 d_1(V) + \omega_2 d_2(V) + \frac{\Lambda}{\delta(1+\underline{r})^2} V^{\psi-1} + \omega_S \underline{r} V$$

And the marginal cost

$$MC(V) = \frac{\omega_1}{B_1^{1/\rho}} + \frac{\omega_2}{B_2^{1/\rho}} + (\psi-1) \frac{\Lambda}{\delta(1+\underline{r})^2} V^{\psi-2} + \omega_S \underline{r}$$

Once again the optimal V is found by equating marginal cost and benefit, which again lacks a closed form solution.

A.2.1 GDPR Effect

There are three outcome changes driven by the GDPR we are interested in: the change in the expected number of breaches, the change in the expected amount of data available in the market, and the change in the average breach size. These will be affected by both the increase in data collection costs and the security constraint, but for the sake of simplicity I will address each separately.

Corollary A.0.1. *Assuming each target has an interior maximum and the distribution of A has a sufficiently high upper bound, there exists a unique $A = A^\dagger$ above which the security ratio constraint binds.*

Proof. Let

$$F(V, A) = MB(V) - MC(V) = 0$$

by the implicit function theorem:

$$\begin{aligned}\frac{\partial V}{\partial A} &= -\frac{\frac{\partial F}{\partial A}}{\frac{\partial F}{\partial V}} \\ \frac{\partial F}{\partial A} &= \phi V^{\phi-1} \\ \frac{\partial F}{\partial V} &= \phi(\phi-1)AV^{\phi-2} - \frac{\partial MC}{\partial V}\end{aligned}$$

$\frac{\partial F}{\partial A}$ is obviously positive, and $\frac{\partial F}{\partial V}$ is the second order condition of the target's objective function. If there is an interior maximum, this will be negative. Therefore, $\frac{\partial V}{\partial A} > 0$.

With that,

$$\frac{\partial r}{\partial A} = \left[\frac{2\Lambda}{\delta\omega_S} \right]^{1/3} \frac{\psi-2}{3} V^{\frac{\psi-5}{3}} \frac{\partial V}{\partial A}$$

Having assumed that $\psi < 2$, r is monotonically decreasing in A . □

Whether π is increasing or decreasing in r depends on r . If $r < 1$, π is increasing in r , while if $r > 1$, π is decreasing. Since r is monotonically decreasing in A , $r < 1$ will be true for targets with sufficiently high A . A property of π is that $\pi(r) = \pi(\frac{1}{r})$:

$$\begin{aligned}\pi(r) &= \frac{r}{(1+r)^2} \\ \pi\left(\frac{1}{r}\right) &= \frac{1/r}{(1+1/r)^2} = \frac{1/r}{\left(\frac{r+1}{r}\right)^2} \frac{r}{(1+r)^2} = \pi(r)\end{aligned}$$

Defining $A = A^{\dagger\dagger}$ as the point where $\pi(r(A^{\dagger\dagger})) = \pi\left(\frac{1}{r}\right)$, if there is a sufficiently small mass of firms where $A > A^{\dagger\dagger}$, then the security requirement constraint will decrease the expected number of breaches.

Any decrease in the expected number of breaches is driven by a decrease in breaches of targets whose A falls between $A^{\dagger}$ and $A^{\dagger\dagger}$. This changes the composition of the data breaches. High A targets will take up a higher portion of breaches than prior to the GDPR. This compositional change will also impact the expected amount of data stolen.

The expected amount of breached data appearing in the market depends on both π and $D(A)$:

$$\mathbb{E}[Data] = \int \pi(A)D(A)dF(A)$$

I have already shown that whether the security ratio constraint causes $\pi(A)$ to increase or decrease will depend on whether A is greater than $A^{\dagger\dagger}$. Total data collection, D , is monotonically

increasing in V , and therefore A . Requiring a higher r increases the marginal cost of data collection, reducing optimal data collection.

Whether the security constraint increases or decreases the amount of data available will depend on the relative sizes of the aforementioned compositional change and the reduction in data collection. The shift towards more data-rich targets increases the expected amount of data, while the decrease in data collection reduces expected data collection. The net change will depend on which of these effects dominates, but the latter has some nuance. Each type of target reduces its data collection, including the now more prominent high A types. But if their reduction in data collection is small enough, they may still have more data than the mid- A types had prior to the GDPR. If this is the case, the composition effect will still increase the expected amount of data stolen.

This same logic applies to determining how the expected breach size changes. The composition effect can still lead to larger breaches on average.

In the stolen data market, consider the price legitimate sellers can expect to receive from a given data package and how that price changes with r :

$$p^* = \frac{S^2V}{(S+V)^2} = \frac{r^2V}{(1+r)^2}$$

$$\frac{\partial p}{\partial r} = \frac{2rV(1+r) - 2r^2V(1+r)}{(1+r)^4} = \frac{2rV}{(1+r)^3}$$

For $r > 0$, price is increasing in r . Price increases in r because q , the portion of legitimate sellers, is also increasing in r :

$$q = \frac{S}{S+V} = \frac{r}{1+r}$$

$$\frac{\partial q}{\partial r} = \frac{1}{(1+r)^2}$$

As r increases, the portion of sellers that are legitimate increases. This increases demand for the good, raising prices. Effectively, the GDPR alleviates the market's lemons problem for large data packages enough to bring more into the market.

References

- Acemoglu, Daron et al. (Nov. 2022). “Too Much Data: Prices and Inefficiencies in Data Markets”. en. In: *American Economic Journal: Microeconomics* 14.4, pp. 218–256. ISSN: 1945-7669. DOI: [10.1257/mic.20200200](https://doi.org/10.1257/mic.20200200). URL: <https://www.aeaweb.org/articles?id=10.1257/mic.20200200> (visited on 12/14/2022).
- Akerlof, George A. (1970). “The Market for “Lemons”: Quality Uncertainty and the Market Mechanism”. In: *The Quarterly Journal of Economics* 84.3, pp. 488–500. ISSN: 0033-5533. DOI: [10.2307/1879431](https://doi.org/10.2307/1879431).
- Aridor, Guy, Yeon-Koo Che, and Tobias Salz (July 2021). “The Effect of Privacy Regulation on the Data Industry: Empirical Evidence from GDPR”. In: *Proceedings of the 22nd ACM Conference on Economics and Computation*. EC ’21. New York, NY, USA: Association for Computing Machinery, pp. 93–94. ISBN: 978-1-4503-8554-1. DOI: [10.1145/3465456.3467543](https://doi.org/10.1145/3465456.3467543). URL: <https://dl.acm.org/doi/10.1145/3465456.3467543>.
- Ayres, Ian and Steven D. Levitt (1998). “Measuring Positive Externalities from Unobservable Victim Precaution: An Empirical Analysis of Lojack”. In: *The Quarterly Journal of Economics* 113.1, pp. 43–77. ISSN: 0033-5533.
- Becker, Gary S. (1968). “Crime and Punishment: An Economic Approach”. In: *Journal of Political Economy* 76.2, pp. 169–217. ISSN: 0022-3808.
- Braakmann, Nils, Arnaud Chevalier, and Tanya Wilson (Oct. 2024). “Expected Returns to Crime and Crime Location”. en. In: *American Economic Journal: Applied Economics* 16.4, pp. 144–160. ISSN: 1945-7782. DOI: [10.1257/app.20220585](https://doi.org/10.1257/app.20220585). URL: <https://www.aeaweb.org/articles?id=10.1257/app.20220585> (visited on 09/26/2024).
- Chen, Chinchih, C. Frey, and Giorgio Presidente (2022). “Privacy Regulation and Firm Performance: Estimating the GDPR Effect Globally*”. In: URL: <https://www.semanticscholar.org/paper/Privacy-Regulation-and-Firm-Performance%3A-Estimating-Chen-Frey/4fb345b55993196739280272a353c6427fdc9f88>.
- Chen, Jiafeng and Jonathan Roth (Dec. 2023). “Logs with Zeros? Some Problems and Solutions*”. In: *The Quarterly Journal of Economics* 139.2, pp. 891–936. ISSN: 0033-5533. DOI: [10.1093/qje/qjad054](https://doi.org/10.1093/qje/qjad054).
- Choi, Jay Pil, Doh-Shin Jeon, and Byung-Cheol Kim (May 2019). “Privacy and personal data collection with information externalities”. In: *Journal of Public Economics* 173, pp. 113–124. ISSN: 0047-2727. DOI: [10.1016/j.jpubeco.2019.02.001](https://doi.org/10.1016/j.jpubeco.2019.02.001). URL: <https://www.sciencedirect.com/science/article/pii/S0047272719300131> (visited on 09/27/2023).
- Deloitte (2020). *Milliseconds Make Millions*. Tech. rep. Deloitte. URL: <https://www.deloitte.com/ie/en/services/consulting/research/milliseconds-make-millions.html>.

- Demirer, Mert et al. (Feb. 2024). “Data, Privacy Laws and Firm Production: Evidence from the GDPR”. In: Working Paper Series 32146. DOI: 10.3386/w32146. DOI: [10.3386/w32146](https://doi.org/10.3386/w32146). URL: <https://www.nber.org/papers/w32146>.
- Department for Digital, Culture, Media and Sport (2022). *Cyber Security Breaches Survey: Combined Dataset, 2016-2022*. data collection. SN: 8971, DOI: <http://doi.org/10.5255/UKDA-SN-8971-1>.
- Franklin, Jason et al. (Oct. 2007). “An inquiry into the nature and causes of the wealth of internet miscreants”. In: *Proceedings of the 14th ACM conference on Computer and communications security*. CCS ’07. New York, NY, USA: Association for Computing Machinery, pp. 375–388. ISBN: 978-1-59593-703-2. DOI: [10.1145/1315245.1315292](https://doi.org/10.1145/1315245.1315292). URL: <https://doi.org/10.1145/1315245.1315292>.
- Goldberg, Samuel G., Garrett A. Johnson, and Scott K. Shriver (Feb. 2024). “Regulating Privacy Online: An Economic Evaluation of the GDPR”. en. In: *American Economic Journal: Economic Policy* 16.1, pp. 325–358. ISSN: 1945-7731. DOI: [10.1257/pol.20210309](https://doi.org/10.1257/pol.20210309).
- Goldfarb, Avi and Catherine Tucker (Jan. 2011). “Privacy Regulation and Online Advertising”. In: *Management Science* 57.1, pp. 57–71. ISSN: 0025-1909. DOI: [10.1287/mnsc.1100.1246](https://doi.org/10.1287/mnsc.1100.1246).
- Herley, Cormac and Dinei Florêncio (2010). “Nobody Sells Gold for the Price of Silver: Dishonesty, Uncertainty and the Underground Economy”. en. In: *Economics of Information Security and Privacy*. Ed. by Tyler Moore, David Pym, and Christos Ioannidis. Boston, MA: Springer US, pp. 33–53. ISBN: 978-1-4419-6967-5. DOI: [10.1007/978-1-4419-6967-5_3](https://doi.org/10.1007/978-1-4419-6967-5_3).
- Holt, Thomas J. and Eric Lampke (Mar. 2010). “Exploring stolen data markets online: products and market forces”. In: *Criminal Justice Studies* 23.1, pp. 33–50. ISSN: 1478-601X. DOI: [10.1080/14786011003634415](https://doi.org/10.1080/14786011003634415).
- Holt, Thomas J., Olga Smirnova, and Yi Ting Chua (Apr. 2016). “Exploring and Estimating the Revenues and Profits of Participants in Stolen Data Markets”. In: *Deviant Behavior* 37.4, pp. 353–367. ISSN: 0163-9625. DOI: [10.1080/01639625.2015.1026766](https://doi.org/10.1080/01639625.2015.1026766).
- Janßen, Rebecca et al. (May 2022). “GDPR and the Lost Generation of Innovative Apps”. In: Working Paper Series 30028. DOI: 10.3386/w30028. DOI: [10.3386/w30028](https://doi.org/10.3386/w30028). URL: <https://www.nber.org/papers/w30028>.
- Jia, Jian, Ginger Zhe Jin, and Liad Wagman (July 2021). “The Short-Run Effects of the General Data Protection Regulation on Technology Venture Investment”. In: *Marketing Science* 40.4, pp. 661–684. ISSN: 0732-2399. DOI: [10.1287/mksc.2020.1271](https://doi.org/10.1287/mksc.2020.1271).
- Jia et al. (June 2025). *How Does Privacy Regulation Affect Transatlantic Venture Investment? Evidence from GDPR*. Working Paper. DOI: [10.3386/w33909](https://doi.org/10.3386/w33909). URL: <https://www.nber.org/papers/w33909> (visited on 03/13/2026).
- Johnson, Garrett A., Scott K. Shriver, and Samuel G. Goldberg (Mar. 2023). “Privacy and Market Concentration: Intended and Unintended Consequences of the GDPR”. In: *Management Science*. ISSN: 0025-1909. DOI: [10.1287/mnsc.2023.4709](https://doi.org/10.1287/mnsc.2023.4709). URL: <https://pubsonline.informs.org/doi/10.1287/mnsc.2023.4709>.

- Jones, Charles I. and Christopher Tonetti (Sept. 2020). “Nonrivalry and the Economics of Data”. en. In: *American Economic Review* 110.9, pp. 2819–2858. ISSN: 0002-8282. DOI: [10.1257/aer.20191330](https://doi.org/10.1257/aer.20191330).
- Kircher, Tobias and Jens Foerderer (Dec. 2021). “Does EU-Consumer Privacy Harm Financing of US-App-Startups? Within-US Evidence of Cross-EU-Effects”. en. In: 4058437. URL: <https://papers.ssrn.com/abstract=4058437>.
- Koski, Heli and Nelli Valmari (2020). *Short-term Impacts of the GDPR on Firm Performance*. eng. 77. URL: <https://www.econstor.eu/handle/10419/237362>.
- Lukic, Karlo, Klaus M. Miller, and Bernd Skiera (Mar. 2023). “The Impact of the General Data Protection Regulation (GDPR) on Online Tracking”. en. In: 4399388. DOI: [10.2139/ssrn.4399388](https://doi.org/10.2139/ssrn.4399388). URL: <https://papers.ssrn.com/abstract=4399388>.
- Miller, Amalia and Catherine Tucker (July 2009). “Privacy Protection and Technology Diffusion: The Case of Electronic Medical Records”. In: *Management Science* 55.7, pp. 1077–1093. ISSN: 0025-1909. DOI: [10.1287/mnsc.1090.1014](https://doi.org/10.1287/mnsc.1090.1014).
- (2011). “Encryption and the loss of patient data”. en. In: *Journal of Policy Analysis and Management* 30.3, pp. 534–556. ISSN: 1520-6688. DOI: [10.1002/pam.20590](https://doi.org/10.1002/pam.20590).
- (Oct. 2018). “Privacy Protection, Personalized Medicine, and Genetic Testing”. In: *Management Science* 64.10, pp. 4648–4668. ISSN: 0025-1909. DOI: [10.1287/mnsc.2017.2858](https://doi.org/10.1287/mnsc.2017.2858).
- Mullahy, John and Edward C. Norton (2024). “Why Transform Y? The Pitfalls of Transformed Regressions with a Mass at Zero”. en. In: *Oxford Bulletin of Economics and Statistics* 86.2, pp. 417–447. ISSN: 1468-0084. DOI: [10.1111/obes.12583](https://doi.org/10.1111/obes.12583).
- Neprash, Hannah, Claire McGlave, and Sayeh Nikpay (Feb. 2026). “Hacked to Pieces? The Effects of Ransomware Attacks on Hospitals and Patients”. en. In: *American Economic Journal: Economic Policy* 18.1, pp. 256–281. ISSN: 1945-7731. DOI: [10.1257/pol.20240594](https://doi.org/10.1257/pol.20240594). (Visited on 01/30/2026).
- Romanosky, Sasha (Dec. 2016). “Examining the costs and causes of cyber incidents”. In: *Journal of Cybersecurity* 2.2, pp. 121–135. ISSN: 2057-2085. DOI: [10.1093/cybsec/tyw001](https://doi.org/10.1093/cybsec/tyw001). URL: <https://doi.org/10.1093/cybsec/tyw001> (visited on 09/08/2023).
- Romanosky, Sasha, Rahul Telang, and Alessandro Acquisti (2011). “Do data breach disclosure laws reduce identity theft?” en. In: *Journal of Policy Analysis and Management* 30.2. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/pam.20567>, pp. 256–286. ISSN: 1520-6688. DOI: [10.1002/pam.20567](https://doi.org/10.1002/pam.20567). URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/pam.20567> (visited on 09/08/2023).
- SpyCloud (2026). *SpyCloud Annual Identity Exposure Report 2026*. Tech. rep. SpyCloud. URL: <https://engage.spycloud.com/rs/713-WIP-737/images/spycloud-2026-identity-exposure-report.pdf>.
- Vu, Anh V. et al. (Oct. 2020). “Turning Up the Dial: the Evolution of a Cybercrime Market Through Set-up, Stable, and Covid-19 Eras”. In: *Proceedings of the ACM Internet Measurement Conference*. IMC ’20. New York, NY, USA: Association for Computing Machinery, pp. 551–566.

ISBN: 978-1-4503-8138-3. DOI: [10.1145/3419394.3423636](https://doi.org/10.1145/3419394.3423636). URL: <https://dl.acm.org/doi/10.1145/3419394.3423636>.

Wooldridge, Jeffrey M (Sept. 2023). “Simple approaches to nonlinear difference-in-differences with panel data”. In: *The Econometrics Journal* 26.3, pp. C31–C66. ISSN: 1368-4221. DOI: [10.1093/ectj/utad016](https://doi.org/10.1093/ectj/utad016).